

OpenStreetMap Data Wrangling with MongoDB

Project

Dmitry Tolstonogov

Problems Encountered in the Map

Map Area: New York, NY United States.

OSM XML file `new-york_new-york.osm` is obtained from [MapZen](#) . We created a smaller sample (1%) of the initial file `sample_1_pct 1.osm` for the exploratory analysis.

We would like to change the data model and expand the "addr:street" type of keys to a dictionary like this:

```
{"address": {"street": "Some value"}}
```

So, we have to see if we have such tags, and if we have any tags with problematic characters.

We modified the script `tags.py` from the problem Tag Types, Lesson 6 of Data Wrangling with MongoDB course in the following way:

For the "k" value for each "<tag>" it checks if they can be valid keys in MongoDB, as well as check if there are any other potential problems.

It uses regular expressions to look for certain patterns in the tags.

The script returns the number of patterns and sets of samples for every pattern.

Using `explore_tags.py` we have found the following problems in the data:

1. There are multiple keys with the prefix like "prefix:", for example:
 - a. `<tag k="addr:street" .../>`
 - b. `<tag k="gnis:state_id" ... />`
 - c. `<tag k="tiger:county" ... />`

- d. `<tag k="railway:atc" ... />`
 - e. `<tag k="contact:fax" ... />`
 - f. `<tag k="source:hgv" ... />`
 - g. `<tag k="building:colour" ... />.`
2. There is [cityrack](#) information with dot separation, like
 - a. `<tag k="cityracks.street" ... />;`
 - b. We should change it similar to the before cases.
 3. There are multiple references in `way` tags.
 4. Some postal codes are incorrect; some are not belongs to New York City.
 5. Some street names are inconsistent.

To deal with the problem 1, we changed the data model and expanded

"some_dict:some_key" type of keys to a dictionary like this: {"some_dict": {"some_key": "some_value"}}. We did it for the all interesting cases discovered in our tag exploration.

Similarly, we did it for the problem 2 (cityrack).

To avoid multiple refs in "way" nodes like this:

```
<nd ref="305896090"/>
<nd ref="1719825889"/>
```

(problem 3) we turned into

```
"node_refs": ["305896090", "1719825889"]
```

The correct postcodes (problem 4) must have 5 digits and belong to the NYC area,

```
10001 <= postcode <= 11692.
```

Also, we corrected cases like `v = 'NY 10533'` or `v = '11229-8541'`

To deal with problem 5 (street name inconsistency), we modified the script from `audit.py` from the problem Improving Street Names, Lesson 6 of Data Wrangling with MongoDB course.

We join all the solutions above to `process_data.py` script. It parses `new-york_new-york.osm` file to `new-york_new-york.osm.json` file.

All the scripts passed the small sampled `sample_1_pct 1.osm` smoothly.

While parsing the big `new-york_new-york.osm` file, we have encountered some challenges.

First, while updating street names, we've got an exception in some place:

```
UnicodeEncodeError: 'ascii' codec can't encode character u'\xf1' in
position 8: ordinal not in range(128). We resolved it applying the method
.encode('ascii', 'ignore').
```

Second, we've got `MemoryError`. We got rid of it clearing each element after its shaping with `element.clear()` command.

Overview of the data

This section contains basic statistics about the dataset.

- **File sizes**

```
sample_1_pct 1.osm is 21162 KB (about 300000 lines)
```

```
sample_1_pct 1.osm is 23159 KB
```

```
new-york_new-york.osm is 2094094 KB (about 30000000 lines).
```

```
new-york_new-york.osm.json is 2315091 KB
```

- **Number of documents**

```
db.count()
```

```
Out[1]: 9831664
```

- **Number of unique users**

```
len(db.distinct("created.user"))
```

```
Out[2]: 2737
```

- **Number of nodes**

```
db.find({"type": "node"}).count()
```

```
Out[3]: 8439186
```

- **Number of ways**

```
db.find({"type": "way"}).count()
```

```
Out[4]: 1392478
```

- **Number of amenities**

```
db.find({"amenity": {"$exists" : 1}}).count()
```

```
Out[5]: 30505
```

- **Number of bridges**

```
len(db.distinct("bridge"))
```

```
Out[6]: 10
```

Indeed, NYC has 10 bridges.

Other ideas about the datasets

- **10 Top amenities**

```
pipeline = [{"$match" : {"amenity" : {"$exists" : 1}}},  
{"$group" : {"_id" : "$amenity", "count": {"$sum" : 1}}},  
{"$sort" : {"count" : -1}}, {"$limit" : 10}]
```

```
res = db.aggregate(pipeline)
```

```
for item in res:
```

```
    pprint.pprint(item)
```

```
{u'_id': u'parking', u'count': 5135}
```

```
{u'_id': u'bicycle_parking', u'count': 4820}
```

```
{u'_id': u'school', u'count': 4628}
```

```
{u'_id': u'place_of_worship', u'count': 4491}
```

```
{u'_id': u'restaurant', u'count': 2280}
```

```
{u'_id': u'fast_food', u'count': 736}
```

```
{u'_id': u'cafe', u'count': 684}
```

```
{u'_id': u'fire_station', u'count': 667}
```

```
{u'_id': u'bank', u'count': 628}
```

```
{u'_id': u'bench', u'count': 477}
```

- **Starbucks cafes distribution by borough**

```
pipeline = [{"$match" : {"name" : {"$regex" : "Starbucks"}} },  
\
```

```

{"$group" : {"_id" : "$address.city", "count": {"$sum" :1}}}, \
{"$sort" : {"count" : -1}}]
res = db.aggregate(pipeline)
for item in res:
    pprint.pprint(item)
{u'_id': None, u'count': 108}
{u'_id': u'New York', u'count': 38}
{u'_id': u'Brooklyn', u'count': 4}
{u'_id': u'Oakland', u'count': 1}
{u'_id': u'brooklyn', u'count': 1}
{u'_id': u'New York City', u'count': 1}
{u'_id': u'Flushing', u'count': 1}
{u'_id': u'Forest Hills', u'count': 1}
{u'_id': u'Ridgewood', u'count': 1}
{u'_id': u'Astoria', u'count': 1}
{u'_id': u'Cos Cob', u'count': 1}
{u'_id': u'Denville', u'count': 1}

```

Obviously, the number of Starbucks cafes in NYC is much larger. The result also outlines the inconsistency of addresses and a room for the further wrangling and improving of OpenStreetMap data.